

Understood. Below is a clean, **high-signal, very detailed** technical write-up documenting exactly what you built, how it worked, how we debugged, and what the root cause was. I am writing it like internal engineering notes rather than tutorial fluff.

---

## Kubernetes + Minikube Flask + MongoDB Full Debug + Deployment Documentation

### 1. Goal / Architecture

You built a local Kubernetes deployment using Minikube with the following components:

- **Flask App**
  - Exposed externally via **NodePort**
  - Connects internally to MongoDB using **ClusterIP service**
- **MongoDB**
  - Runs as a **Deployment**
  - Has a **PersistentVolume (PV)** and **PersistentVolumeClaim (PVC)** for data durability
  - Auth enabled via env variables
- **Storage**
  - PV: hostPath
  - PVC: bound to PV
- **Networking**
  - Internal: service name mongodb-service
  - External: NodePort on Minikube

Architecture Diagram:

Browser → NodePort Service → Flask Pods → MongoDB Service → MongoDB Pod → PV/PVC

---

### 2. Files / Manifests You Used

You created the following YAMLs:

#### (a) Flask Deployment

- 3 replicas
- readiness & liveness probes
- env vars for Mongo credentials

#### (b) Flask Service (NodePort)

Used to expose Flask publicly:

type: NodePort

```
port: 5000  
nodePort: 30000
```

#### (c) MongoDB Deployment

- Uses containerPort: 27017
- Uses MONGO\_INITDB\_ROOT\_USERNAME and MONGO\_INITDB\_ROOT\_PASSWORD
- Mounts volume /data/db

#### (d) MongoDB Service (ClusterIP)

Internal service, DNS name becomes:

```
mongodb-service.default.svc.cluster.local
```

#### (e) PersistentVolume + PersistentVolumeClaim

PV used hostPath:

```
hostPath: mongo/data
```

PVC requested 2Gi storage.

---

### 3. Minikube Execution Workflow (Correct Order)

This is the correct minimal workflow you followed:

#### Step 1: Start Minikube

From project root:

```
minikube start
```

#### Step 2: Configure Docker to Build Images Inside Minikube

Critical step for local images to be visible to Kubernetes:

Windows (CMD):

```
FOR /f "tokens=*" %i IN ('minikube -p minikube docker-env') DO %i
```

This sets environment variables:

```
DOCKER_HOST
```

```
DOCKER_TLS_VERIFY
```

```
MINIKUBE_ACTIVE_DOCKERD
```

#### Step 3: Build Flask Image

From project root:

```
docker build -t flask-app .
```

This ensures that the Kubernetes cluster can pull flask-app locally (no registry push needed).

## **Step 4: Apply Kubernetes Manifests**

Move into manifests folder:

```
cd Manifests
```

```
kubectl apply -f .
```

This created:

- Deployments
- Services
- PV & PVC

## **Step 5: Verify Resources**

Commands used:

```
kubectl get pods
```

```
kubectl get svc
```

```
kubectl get pv
```

```
kubectl get pvc
```

## **Step 6: Access Flask**

Option A:

```
minikube service flask-service --url
```

Option B (manual):

```
minikube ip
```

```
kubectl get svc flask-service
```

Then browse:

```
http://<minikube-ip>:30000
```

---

## **4. Debugging Session Recap (Very Detailed)**

### **Issue #1: CrashLoopBackoff on MongoDB**

Initial logs showed:

```
DBPathInUse: Another mongod instance is already running
```

Meaning:

- Old MongoDB pod **already owned /data/db**
- Your new MongoDB tried to use same volume → lock conflict

**Root Cause:** Old Stateful workloads still using the PV/PVC from 309 days ago.

#### Action Taken:

1. Examined running pods:

```
kubectl get pods
```

Saw old deployments: flask-todo, mongodb, etc.

2. Checked PVC:

```
kubectl get pvc
```

Saw mongo-pvc in Terminating state bound to old PV.

3. Deleted old RC and Deployments:

```
kubectl delete deployment flask-todo ...
```

```
kubectl delete rc todo-rc
```

4. Deleted legacy PV/PVC:

```
kubectl delete pvc mongo-pvc
```

```
kubectl delete pv <pv-name>
```

5. Re-applied new PV/PVC

```
kubectl apply -f pvc.yml
```

```
kubectl apply -f pv.yml
```

Result:

- MongoDB could now start cleanly
- Lock file issues disappeared

#### Conclusion of Debug #1

Mongo **never had authentication or image issues**.

The root cause was **old pods occupying the same PVC path**, not credentials.

---

#### Issue #2: Flask CrashLoopBackoff

Initially, Flask pod crashed because Mongo pod wasn't ready.

Logs showed socket connection refused.

Sequence:

- Flask started → tried Mongo connect → Mongo not up → container crash
- Kubernetes restarts container → eventually all stabilize

After fixing Mongo:

```
kubectl get pods
```

Showed all Flask pods READY = 1/1

---

## 5. Networking Resolution & Verification

Internal resolution:

Flask env variables:

MONGO\_HOST=mongodb-service

MONGO\_PORT=27017

Inside cluster:

mongodb-service.default.svc.cluster.local

External resolution:

minikube ip → 192.168.49.2

NodePort → 30000

Final browser URL:

<http://192.168.49.2:30000>

or via helper:

minikube service flask-service --url

This launched browser → Flask served HTML → Confirmed DB ops work.

---

## 6. Authentication Clarification

On Docker Compose:

- Mongo auth succeeded

On Kubernetes:

- Auth also succeeded once Mongo was booted correctly
- No change was actually required to authentication logic
- The real blockers were stale PVC + stale pods, not credentials

You correctly concluded here:

"Most of my issue wasn't due to any authentication error. It was mainly because my earlier pods were running and they were messing up."

That is technically accurate.

---

## 7. Key Kubernetes Concepts You Proved

You successfully demonstrated:

Concept	Verified By
Cluster internal DNS	Flask → mongodb-service
Service discovery	selector + endpoints
Pod restart policy	CrashLoopBackoff recovery
Persistent storage	PVC + PV binding
Docker build inside cluster docker-env env switch	
External access	NodePort via minikube ip
Debugging logs	kubectl logs
Resource cleanup	deleting old RC/PVC/PV

This is basically a full production-grade workflow minus secrets.

---

## 8. Lessons Learned / Retrospective

1. **Minikube still retains old workloads** from previous runs unless explicitly deleted
  2. **PVCs survive pod deletion**, so stale locks persist
  3. **DB containers fail hard** on data directory locks
  4. **NodePort vs ClusterIP is contextual**
  5. **CrashLoopBackOff is often readiness not auth**
  6. **Local Docker images require docker-env** to be recognized by Minikube
- 

## 9. Final Working State

At the end:

```
kubectl get pods
```

showed:

```
flask-app-deployment XXX 1/1 Running
```

```
mongodb-deployment XXX 1/1 Running
```

and:

```
minikube service flask-service --url
```

opened the app successfully.

---

## **10. Optional Improvements (Future)**

Not necessary, but recommended for production:

- Use StatefulSet for MongoDB
  - Use Secrets instead of raw env vars
  - Use ConfigMap for non-secret configs
  - Use LoadBalancer on EKS instead of NodePort
  - Use StorageClass for dynamic provisioning
- 

If you want, I can also convert this into:

- Markdown doc
- PDF
- Notion wiki page
- GitHub README
- Technical blog post format
- Interview case study format

Just reply:

markdown, pdf, notion, readme, blog, or interview.

**Yes, that is the correct high-level flow**, and your summary is almost perfect. I'll restate it cleanly with the precise ordering and reasoning so there's zero ambiguity:

---

### **Correct End-to-End Flow (Clean Summary)**

#### **STEP 1 — Author Kubernetes YAML Manifests**

You create Kubernetes manifests for:

1. **MongoDB Deployment**
2. **MongoDB Service (ClusterIP)**
3. **Flask Deployment**
4. **Flask Service (NodePort for external access)**
5. **PersistentVolume (PV)**
6. **PersistentVolumeClaim (PVC)**

At this point you have **desired state definitions**, not running resources.

---

## STEP 2 — Start Kubernetes Cluster (Minikube)

`minikube start`

This boots a single-node Kubernetes cluster locally.

---

## STEP 3 — Build Application Docker Image INSIDE Minikube Docker Daemon

Reason:

- Kubernetes pulls images from registries
- Your image is local-only
- So Minikube must “see” it

Switch Docker context:

Windows CMD:

```
FOR /f "tokens=*" %i IN ('minikube -p minikube docker-env') DO %i
```

Then build:

```
docker build -t flask-app .
```

The crucial point:

**The name flask-app must match the image name in the Deployment YAML**, otherwise Kubernetes won’t find it.

This eliminates the need for Docker Hub.

---

## STEP 4 — Apply Kubernetes Manifests

Go to your Manifests folder:

```
kubectl apply -f .
```

Kubernetes now creates:

- Deployments (Mongo + Flask)
- Services (Mongo + Flask)
- PV + PVC

This is where pods start being scheduled.

---

## STEP 5 — Kubernetes Internal Boot Order Happens Automatically

Behind the scenes:

1. **MongoDB pod starts**
  2. Mongo mounts PVC
  3. MongoDB Service (ClusterIP) gets a stable DNS like:
  4. `mongodb-service.default.svc.cluster.local`
  5. Flask pod starts
  6. Flask reads env vars:
  7. `MONGO_HOST=mongodb-service`
  8. Flask connects to DB through service, not by IP  
(so scaling or rescheduling will not break)
  9. Flask Service (NodePort) exposes app externally
- 

## **STEP 6 — Access Application Externally**

Option A (easiest):

```
minikube service flask-service --url
```

Option B (manual):

```
minikube ip
```

```
kubectl get svc flask-service
```

Open:

```
http://<minikube-ip>:<nodePort>
```

---

## **So Your Final Summary is Correct**

Your exact statement reformulated:

“We write our deployments/services for both Flask and Mongo, then we define PV/PVC for storage, then we start Minikube, build the Flask Docker image inside Minikube, apply all manifests, and Kubernetes brings everything up. After that, application runs.”

Yes.

---

## **One Slight Technical Clarification**

When you said:

“so my deployment YAML can find it”

More precisely:

- Kubernetes doesn't directly read from your Docker build command
- Kubernetes asks **container runtime** to pull the image
- Since you built the image inside Minikube's Docker daemon, Kubernetes finds it there

So the real meaning is:

The deployment pulls the image because it exists in the Minikube Docker daemon and matches the name in the Deployment YAML.

---

### Final Flow in One Sentence

**Define manifests → Build image inside Minikube → Apply manifests → Kubernetes creates pods/services → Flask connects to Mongo via service → App becomes externally accessible via NodePort.**

That's it.

---

If you want, I can now give you:

- A condensed cheat sheet, or
- A full "AWS EKS version" of the same workflow (with LoadBalancer + ECR), or
- A full CI/CD GitHub Actions pipeline for automation.

Just say: **cheatsheet**, **EKS version**, or **CI/CD**.